


# Build Your Own Router

Computer Networks  
Course Project

# Project Overview

- Implement a simple router in a single router topology with **static routing table** (forwarding table).
- Your router will receive **raw Ethernet frames**, and handle/forward packets in correct logic.




**LESSONS TO LEARN**

# What about Our Router?

- **Forwarding:** move packets from router's input to appropriate router output
  - Load predefined routing table
  - Look up matching entry in routing table
- Handle ICMP
- Handle ARP request / reply
- ~~Routing: determine route taken by packets from source to dest.~~

# Link Layer


- Transfer internet layer datagram from **node to adjacent node** over a link
- Encapsulate internet layer **datagram into frame**, add header & trailer
- Use "**MAC**" address in frame headers for source, dest



# MAC address v.s. IP address

- 32-bit IP address:
  - network-layer address
  - used to get datagram from src to dest IP subnet
- 48-bit MAC (Ethernet) address:
  - link-layer address for network interfaces
  - get frame from **one interface to another physically-connected interface (same subnet)**
  - Broadcast address: "FF:FF:FF:FF:FF:FF"


# Ethernet Frame



- Encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**
  - preamble and crc are handled transparently in this project


# ARP: Address Resolution Protocol

Question: How to get MAC address of B from B's IP address




- ARP request: request IP-MAC mapping of next hop interface (send to broadcast address)
- ARP reply: send IP-MAC of current interface
- ARP cache: **IP - MAC mapping for nodes** (timeout after a 30s)

# ARP: Address Resolution Protocol



Walkthrough: send datagram from A to B through R.

# Internet Protocol (IPv4)




- Delivering packets from src to dest based on IP address
- Header **checksum**
- Decrement **TTL**

# Routing Table

- Destination & Netmask: subnet **network ID**
- Gateway: next hop IP address to destination
- Interface: name of network interface card connected to gateway
- Metric: routing metric of path to destination  
(omitted in this project)

# Longest Prefix Match


Destination	Netmask	Gateway	Interface
0.0.0.0	0.0.0.0	10.0.1.100	eth3
192.168.2.2	255.255.255.0	*	eth1
172.64.3.10	255.255.0.0	*	eth2



Where should packet dest at "192.168.1.1" go?


# Longest Prefix Match

Destination	Netmask	Gateway	Interface
0.0.0.0	0.0.0.0	10.0.1.100	eth3
192.168.2.2	255.255.255.0	*	eth1
172.64.3.10	255.255.0.0	*	eth2



Where should packet dest at "192.168.2.1" go?

# ICMP: Internet Control Message Protocol



- Used by hosts & routers to communicate network-level information
  - error reporting (unreachable host, network, port, protocol)
  - echo request/reply
- Sent as **IPv4 payload**
- Content: Internet header + 8 bytes of original datagram

# Ping & Traceroute

- Ping: ICMP echo request/reply
  - reply with TTL = 64

# Ping & Traceroute

- Ping: ICMP echo request/reply
  - reply with TTL = 64
- Traceroute: displaying possible routes and RTT in IP network.
  - sends UDP segments with TTL = 1,2,3, ... with unlikely port number
  - when nodes receive datagram with TTL=0, return ICMP "Time Exceeded"
  - destination returns ICMP "Port Unreachable"


# Useful Materials

- IPv4:
  - RFC 791: <https://tools.ietf.org/html/rfc791>
  - Text Book: section 4.4.1, 4.4.2
- ICMP:
  - RFC 792: <https://tools.ietf.org/html/rfc792>
  - Text Book: section 4.4.3
- ARP:
  - RFC 826: <https://tools.ietf.org/html/rfc826>
  - Text Book: section 5.4.1

**RESULTS TO SHOW**

# Expected Behaviors

- ping from client to any server & router interfaces
- traceroute from client to any server & router interfaces
- wget files from server
- Update ARP cache table



# Ping

- ping from client to any http servers

```
mininet> client ping server1
PING 192.168.2.2 (192.168.2.2) 56(84) bytes of data.
64 bytes from 192.168.2.2: icmp_seq=1 ttl=63 time=1293 ms
64 bytes from 192.168.2.2: icmp_seq=2 ttl=63 time=312 ms
64 bytes from 192.168.2.2: icmp_seq=3 ttl=63 time=50.3 ms
64 bytes from 192.168.2.2: icmp_seq=4 ttl=63 time=29.3 ms
^C
--- 192.168.2.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3025ms
rtt min/avg/max/mdev = 29.388/421.625/1293.814/515.809 ms, pipe 2
```

# Ping

- ping wrong IP address

```
mininet> client ping 192.168.2.3
PING 192.168.2.3 (192.168.2.3) 56(84) bytes of data.
From 10.0.1.1 icmp_seq=1 Destination Host Unreachable
From 10.0.1.1 icmp_seq=2 Destination Host Unreachable
From 10.0.1.1 icmp_seq=3 Destination Host Unreachable
From 10.0.1.1 icmp_seq=4 Destination Host Unreachable
From 10.0.1.1 icmp_seq=5 Destination Host Unreachable
^C
--- 192.168.2.3 ping statistics ---
6 packets transmitted, 0 received, +5 errors, 100% packet loss, time 5100ms
pipe 5
```

# Traceroute

- traceroute to any http server

```
mininet> client traceroute server1
traceroute to 192.168.2.2 (192.168.2.2), 30 hops max, 60 byte packets
 1  10.0.1.1 (10.0.1.1)  12.806 ms  13.727 ms  14.505 ms
 2  192.168.2.2 (192.168.2.2)  99.179 ms  104.646 ms  106.050 ms
```

- traceroute to router interfaces

```
mininet> client traceroute 192.168.2.1
traceroute to 192.168.2.1 (192.168.2.1), 30 hops max, 60 byte packets
 1  10.0.1.1 (10.0.1.1)  796.441 ms  839.769 ms  839.814 ms
```

# File Downloading

- wget from any of servers

```
mininet> client wget http://192.168.2.2/tmp
--2021-11-09 17:18:15--  http://192.168.2.2/tmp
Connecting to 192.168.2.2:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 10240 (10K) [application/octet-stream]
Saving to: ‘tmp’

tmp          100%[=====>]  10.00K  --.-KB/s    in 0.06s

2021-11-09 17:18:17 (163 KB/s) - ‘tmp’ saved [10240/10240]
```

# ARP Cache table

- “pingall”


Every 1.0s: ./show-arp.py		Tue Nov 9 17:20:53 2021		
MAC	IP	AGE	VALID	
<hr/>				
9a:45:f8:6e:47:2b	192.168.2.2	2 seconds	1	
52:b1:d0:a7:a9:c9	10.0.1.100	1 seconds	1	
82:10:d6:23:40:ad	172.64.3.10	0 seconds	1	

- 30 seconds later

Every 1.0s: ./show-arp.py		Tue Nov 9 17:21:24 2021		
MAC	IP	AGE	VALID	
<hr/>				

**CODE TO IMPLEMENT**

# Code Structure



# Key Methods

- NEED TO IMPLEMENT
- Method that receives a raw Ethernet frame  
(simple-router.hpp|cpp):

```
/**  
 * This method is called each time the router receives a packet on  
 * the interface. The packet buffer \p packet and the receiving  
 * interface \p inIface are passed in as parameters.  
 */  
void  
SimpleRouter::handlePacket(const Buffer& packet, const std::string& inIface);
```

# Key Methods

- **IMPLEMENTED**
- Method to send raw Ethernet frames (`simple-router.hpp|cpp`):

```
/**  
 * Call this method to send packet \p packt from the router on interface \p outIface  
 */  
void  
SimpleRouter::sendPacket(const Buffer& packet, const std::string& outIface);
```

# Key Methods

- NEED TO IMPLEMENT
- Method to handle ARP cache events (arp-cache.hpp|cpp):

```
/**  
 * This method gets called every second. For each request sent out,  
 * you should keep checking whether to resend a request or remove it.  
 */  
void  
ArpCache::periodicCheckArpRequestsAndCacheEntries();
```

# Key Methods

- NEED TO IMPLEMENT
- Method to lookup entry in the routing table (routing-table.hpp|cpp):

```
/**  
 * This method should Lookup a proper entry in the routing table  
 * using "Longest-prefix match" algorithm  
 */  
RoutingTableEntry  
RoutingTable::lookup(uint32_t ip) const;
```

## **ISSUES TO NOTICE**

# 环境配置

- 作业文档中提供了详细的环境配置指南
- 为了方便同学们快速配置环境，我们也提供了虚拟机文件
  - Windows VMware:  
<https://cloud.tsinghua.edu.cn/d/05b4618dad6046b49438/>
  - Mac(Apple Silicon) UTM:  
<https://cloud.tsinghua.edu.cn/f/d09f8e35213540a5a2cf/>
- 虚拟机的账号和密码都是 **router**

# Some Important Issues

- Grading (up to 105%)
  - Router Implementation (85% = 45% public + 40% private)
    - Ping tests (50%)
    - Traceroute tests (20%)
    - File Downloading tests (15%)
  - Project Report + Code Quality (20%)
- Individual work

# Some Important Issues

- Submission
  - Source code ("make tarball")
  - Do not modify **existing data structures**
  - Report: no longer than **THREE** pages
- Evaluation
  - ping, traceroute, file downloading, details in **project spec**
  - Code quality
  - Project report

# 补交规则

- 正常提交: 在 DDL 前提交作业正常计分
- 最迟提交期限和惩罚
  - 每次作业最迟在 DDL 后一周 (7天) 内提交, 超出此期限一律拒收
  - 未超出此拒收期限的迟交作业得分\*0.8.
- 宽限期: 全部大作业共享 7 天宽限期。迟交累计不超过该期限的免于扣分惩罚

# Some Important Issues

- Start up: today, after class
- Deadline: 12.21, 23:59 (3 weeks)
- Late Submission: 2025.12.22 ~ 2025.12.28, 23:59
  - Score \* 0.8
  - No more submissions after 2025.12.28

# Problems Emerged

- Improper TTL handling in traceroute / ping
  - when to do TTL--?
  - when to send ICMP time exceeded?
- Fail to maintain ARP cache entries / send ARP requests
- Imperfect longest prefix match
- Compilation & project structure problem
  - check before submission

Get to work  
as soon as possible!

# Q&A

Good luck

# Acknowledgement

- This project is based on the CS118 class project by Alexander Afanasyev, UCLA.